


Extracting concepts from
a pretrained SAE model's hidden state



Learning to predict concepts &
Mixing/Interleaving continuous concepts into the hidden state