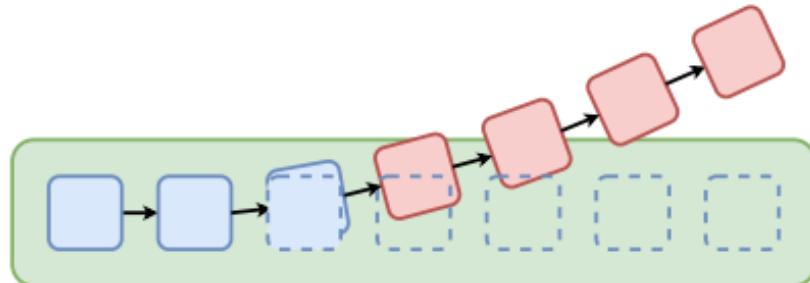
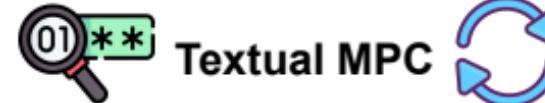
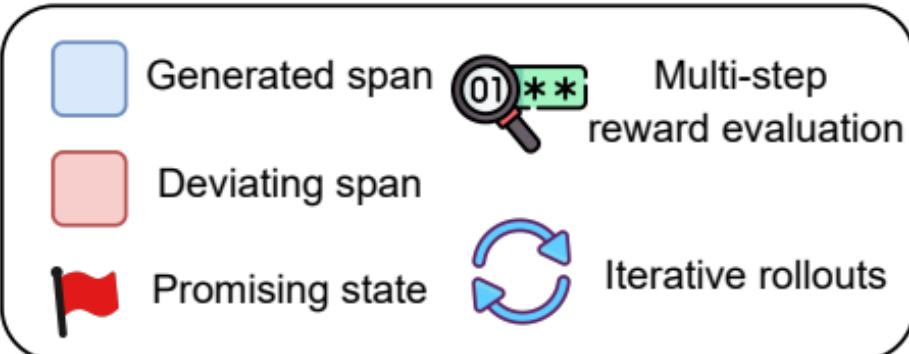




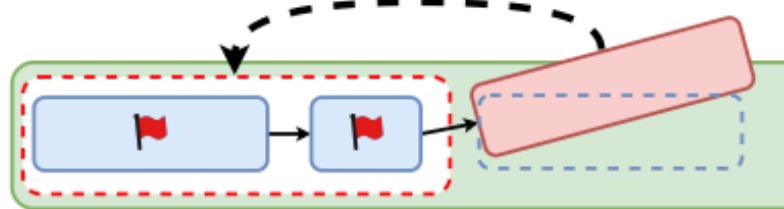
Curse of Horizon (Guided Decoding)



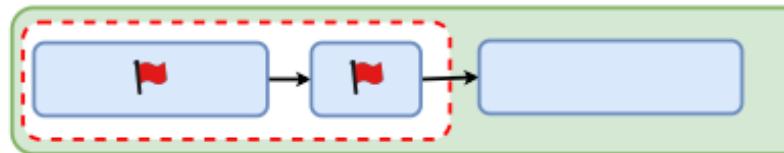
Token-level control leads to overly long horizons, making alignment fragile.



Textual MPC



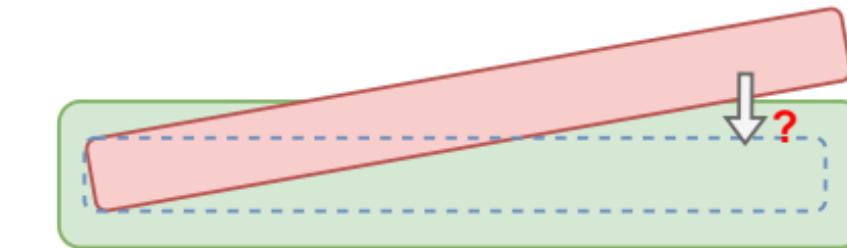
Hindsight Subgoal Identification



Subgoal-Conditioned Re-Generation



Curse of Dimensionality (Iterative Refinement)



Response-level rewriting operates in a large action space, causing unstable improvements.

